

WebNPR

Alan Lee, Jordan Hochman, Maya Diaz Huizar

Introductions



Alan Lee

Master's in Computer Graphics
and Game Technology

LinkedIn: [/in/soohyun-alan-lee/](#)

GitHub handle: [/Alan7996](#)

Personal website: [alannos.com](#)



Jordan Hochman

Master's in Computer Science

LinkedIn: [/in/jhochman24/](#)

GitHub handle: [/JHawk0224](#)

Personal website: [jordanh.xyz](#)



Maya Diaz Huizar

Master's in Computer Science

LinkedIn: [/in/maya-diaz-huizar/](#)

GitHub handle: [/Aorus1](#)

WebGPU Pathtracer + NPR, Why?

Goal: Explore WebGPU as a platform for artists' exploration of stylization and simulation

- WebGPU-based pathtracer
- Stylized Rendering as a Function of Expectation (SIGGRAPH 2024)
- Cloth Simulation



NPR how?


1. On each material shade kernel call...
2. If the current parameters qualify for ANY stylization...
 - a. Generate additional rays to estimate incoming radiance
 - b. Apply stylization function as expectation on these queries

“StyleContext” captures

- ObjectId, MaterialId, Path Prefix, World position, Surface normal, Incoming ray direction, BVH node index, etc.


Physically-based Rendering


$$L_e(\mathbf{x}, \mathbf{y}) + \int_{\mathcal{V}} f_r(\mathbf{x}, \mathbf{y}, \mathbf{z}) G(\mathbf{y}, \mathbf{z}) L(\mathbf{y}, \mathbf{z}) d\mathbf{z}$$





Stylized Rendering

$$g_{\theta} \left(L_e(\mathbf{x}, \mathbf{y}) + \int_{\mathcal{V}} f_r(\mathbf{x}, \mathbf{y}, \mathbf{z}) G(\mathbf{y}, \mathbf{z}) L(\mathbf{y}, \mathbf{z}) d\mathbf{z} \right)$$





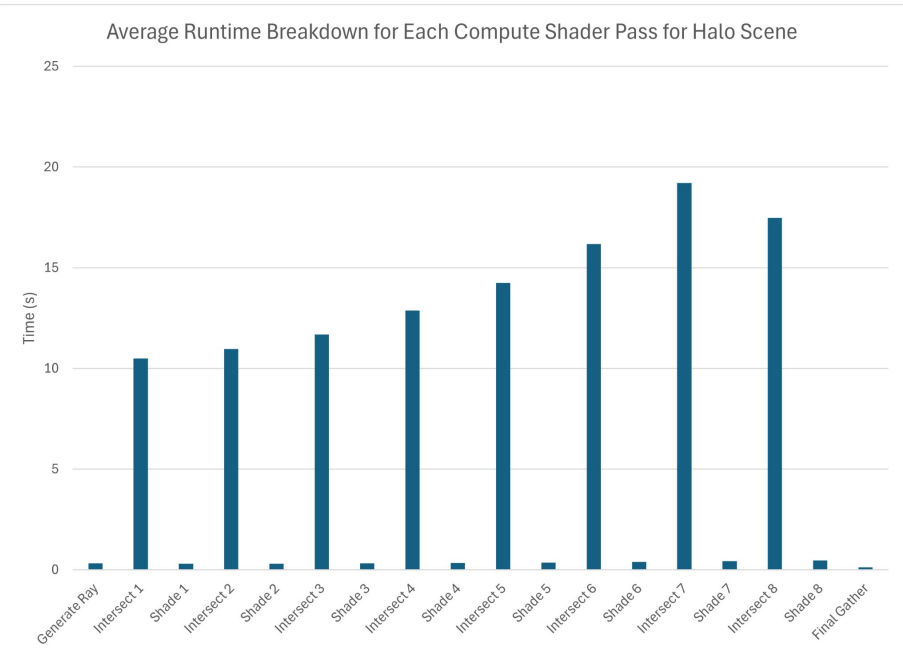

Results






Hero's Room

Performance



- Surprisingly, gained performance WITH NPR
- Very tiny material shade compute pass runtime explains minimal slow down, but why faster?

Thank you!



Alan Lee

alanlee.shl@gmail.com

Jordan Hochman

jhawk111@icloud.com

Maya Diaz Huizar

mayajhuizar@gmail.com

Credits

- Stylized Rendering as a Function of Expectation (2024) (<http://cv.rexwe.st/pdf/srfoe.pdf>)
- Progressive Simulation for Cloth Quasistatics (2023) (<https://pcs-sim.github.io/pcs-main.pdf>)
- Progressive Dynamics for Cloth and Shell Animation (2024) (<https://pcs-sim.github.io/pd/progressive-dynamics-main.pdf>)
- Vite (<https://vitejs.dev/>)
- loaders.gl (<https://loaders.gl/>)
- dat.GUI (<https://github.com/dataarts/dat.gui>)
- stats.js (<https://github.com/mrdoob/stats.js>)
- wgpu-matrix (<https://github.com/greggman/wgpu-matrix>)
- wgpu-basecode (<https://github.com/CIS5650-Fall-2024/Project4-WebGPU-Forward-Plus-and-Clustered-Deferred>)
- Halo Spartan Model (<https://skfb.ly/6QVvM>)
- Suzanne (<https://github.com/KhronosGroup/glTF-Sample-Models/tree/main/2.0/Suzanne>)
- Person (<https://poly.pizza/m/qbDLeTtb8K>)
- Skull (<https://sketchfab.com/3d-models/skull-downloadable-1a9db900738d44298b0bc59f68123393>)
- Small Airplane (<https://poly.pizza/m/7cvx6ex-xfL>)

Credits



From Rex West et al., Stylized Rendering as a
Function of Expectation (2024)

Credits

Algorithm 1: Tree sampling for stylized rendering

```

1 for  $pixel \in image$  do
2    $ray = camera.sampleRay(pixel)$ 
3    $x = ray.origin$ 
4    $y = scene.intersects(ray)$ 
5    $sample = W_p(x, y)G(x, y)stylize(x, y)/p(x, y)$ 
6    $image[pixel].accumulate(sample)$ 

Input:  $x$ : the previous vertex,  $y$ : the current vertex,  $\theta$ : the stylization
parameters (e.g. path prefix, vertex depth, position, time, etc.)

Output:  $L_s$ : the stylized radiance

7 Function  $stylize(x, y)$ :
8    $style = determineStyle(\theta, y)$ 
9    $sampleCount = style.requiredSamples(\theta, y)$ 
10   $samples = []$ 
11  for  $i \in \{0, \dots, sampleCount - 1\}$  do
12    if  $isTerminal(y)$  then
13       $L_o = L_e(x, y)$ 
14    else
15       $ray = sampleNextRay(x, y)$ 
16       $z_i = scene.intersects(ray)$ 
17       $L_i = stylize(y, z_i)$ 
18       $L_o = L_e(x, y) + f_r(x, y, z_i)G(y, z_i)L_i/p(z_i|x, y)$ 
19     $samples.add(L_o)$ 
20   $L_s = style.apply(\theta, samples)$ 
21  return  $L_s$ 

```


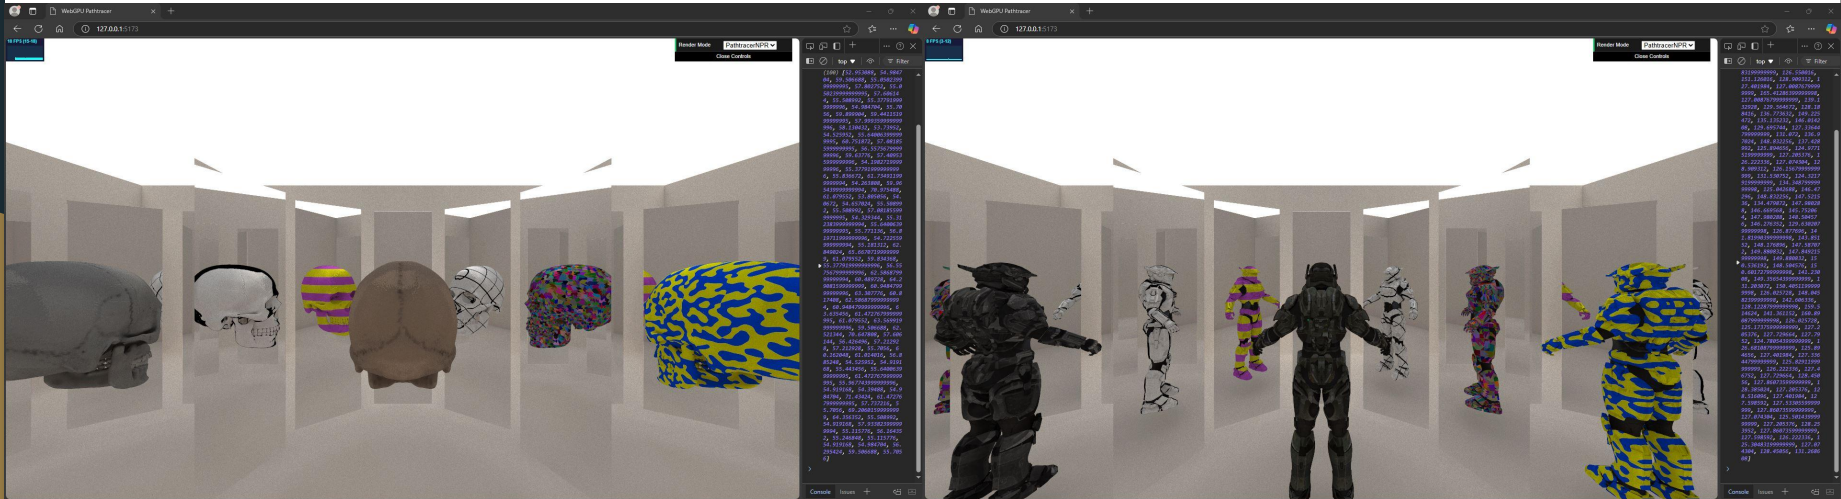



Fig. 6. Evaluating a group-unbiased estimator at a surface point may require more than one integrand sample (e.g. the blue vertex on the bunny requires 3 samples). Starting from the camera and recursively expanding outwards, we observe that satisfying the sample requirements of each next gu-estimator forms a tree (bottom), similar to that of distribution ray tracing [Cook et al. 1984].

From Rex West et al., Stylized Rendering as a Function of Expectation (2024)


Appendix - Test Scene Setup



Appendix - Total Shader Runtimes for Skull




Appendix - Total Shader Runtimes for Halo




Appendix - Shader Runtime Comparisons

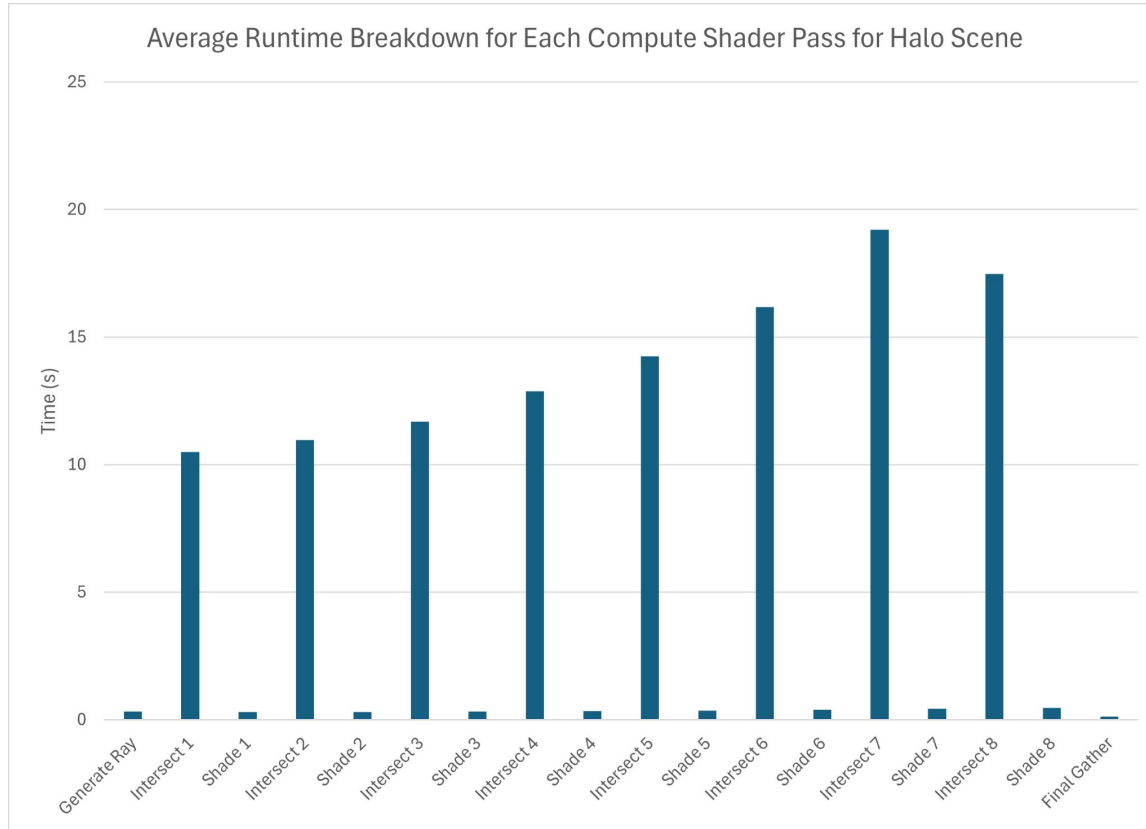
Shader Runtime Comparison for Skull Scene




Shader Runtime Comparison for Halo Scene




Appendix - Compute Shader Pass Runtime Breakdown



Appendix - Cloth Sim Granularities



50 subdivisions per axis



500 subdivisions per axis